
CMPE 239 – Web and Data Mining

Spring 2017

Group Project

Project Timeline

- 3/9: Project proposal DUE
 - Teams formed
 - Dataset selected
 - Submit based on the instructions
- 5/4: Project report & presentation slides due on Canvas
- 5/11 & 5/16: Project presentations

* This is a group project (groups of 3 students). However, each student will be graded individually (please refer to the green sheet for details).

* A list of online resources is posted on Canvas.

* All project documentation (project proposal, project report) should be submitted in the appropriate assignment in Canvas (one per team).

* Provide appropriate citations in your written work, including for datasets. Plagiarism will not be tolerated.

Project proposal

Should be 1-2 pages long and include:

- Student names
- Team name (optional)
- Project title
- Project description
- Proposed methodology/techniques/resources/datasets to be used (see details below)
 - It is important to secure an adequately large dataset before committing to an idea.

General guidelines

The project should be an application that executes an in-depth data mining analysis of a large dataset. You should integrate/combine one or more data mining techniques (possibly using available tools) into a solution for a specific problem.

- The project will be carried out by teams of 3 students. If you are unable to find a team or are missing a member, please use the discussion forum to search for teammates.
- Project deliverables:
 - implementation (source code, binary, and instructions on how to install/use required software)
 - experimental results (evaluation of algorithms, parameters etc.)
 - final report
 - presentation slides
- You may use publicly available (open source) code in your application and existing tools/platforms but this should be explicitly stated in your report. You can choose any programming language to implement your project. You are encouraged to use Python libraries (SciPy, NumPy, Pandas, NLTK, Scrapy, Scikits-Learn), Spark's MLlib, Mahout on AWS, R libraries, or develop your own code.

If you intend to use parts of a previous or existing project from other class (or your MS project), you need to talk to me in order to get approval (not guaranteed).

Choosing a topic

In this project, the team is responsible for proposing a topic to work on. Even though I am providing you with a list of suggestions, it is preferable (and more fun!) to pick a problem “closer to your heart”. Just make sure that data are available that you can use before writing the proposal (no data availability means that you will have to revise your proposal and waste valuable time).

Therefore, it is highly recommended that you look at available data sets that look interesting (and are large enough) and then try to formulate a problem/application around them. Many links to datasets are posted on Canvas. However, you are free to explore other options and/or consider writing a script that will collect data from existing sources (again, you have to ensure that you’ll be able to do this before committing to a project).

All the projects should include one or more (depending on the scope) data mining algorithms.

A few sample projects are (these are provided mainly to give you an idea of the scope of the project):

- Design a program that will classify review comments (on products, blogs, etc.) as spam/fake using sentiment analysis
- Design a recommendation engine that will generate recommendations taking into account other factors (e.g., social network connections, item reviews)
- Design a recommendation system that will bundle items (e.g. a good hotel and a good restaurant when visiting a city)
- Design a recommendation system or perform some data analysis using fashion-inspired data (e.g. combine outfits)
- Design a predictive model for sports events or sport-related activities (e.g. predict which team is going to win the super bowl).
- Design a program that will use/analyze data from social networks. Note that there are so many possible ways of doing this – e.g. analyze tweets in different US states and find the “happiest” folks, then show results on a map.
- Develop your own system that will mine the social graph of a network and generate interesting patterns (e.g. important users).
- Use the data provided by a city (e.g. San Jose or San Francisco) and see what interesting patterns & applications you can implement that would be useful in a community.
- Use data collected by smart devices (e.g. activity, geo-location, etc.) and design an application that helps the user in some way.

Topics to avoid (many others have done these before):

- Movie/Song/Book recommendation system
- Sentiment analysis/Word clouds on Twitter
- Anything using the WEKA GUI (too simple!)

If your group is looking for a challenge, above and beyond the minimum required for this assignment, come talk to me during office hours. I have a couple very interesting projects needing some interested groups to work on them.

Deliverables

Project Report

- Report **should not exceed 10 pages** (additional details if needed can be included in an Appendix. However, the data in the appendix will not contribute to/against your final grade and should be only informative.)
- A project report template has been included in the Appendix. You need to update/adjust the template to your project's focus and details.

Presentation

Presentation slides should be uploaded to Canvas. All team members need to present some part of the talk. Depending on the type of the project, you can include a short demo of your prototype, but the demo counts as part of your presentation and no additional time will be allocated. Each team will have 8-10 minutes to present.

Code & Binary

You will need to submit your project distribution files and documentation, *both electronically and in hard copies* (Canvas upload & printout of the report).

Your project distribution files should contain the following:

1. Complete source code (be sure that your source code is fully documented and easy to read), including data pre-processing scripts.
2. Binary files (e.g., executables, DLLs, Class files) or other components necessary to run your program.
3. Readme file containing instructions on how to compile, install, and/or run your program.
4. Any test data used for evaluation of your system.

Evaluation

The projects will be evaluated based on significance, design, correctness, documentation, and appropriate evaluation/testing. A grading rubric for the project has been included in the Appendix.

Each student will be evaluated individually. Your final grade will be calculated based on:

- The overall quality of the project report
- The overall quality of the project prototype (implementation)
- The overall quality of the experimental evaluation of your approach
- The individual presentation and overall quality of the presentation
- The individual participation in all project aspects (as shown by the project timeline and the peer evaluation)

APPENDIX

Sample of previous CMPE 239 project topics

Please note that this is for your reference, in order to understand the expected scope and extend of work. However, I will not approve projects that cover the exact same topics (combination of idea & dataset) as listed below.

- Movie recommendation system using Matrix Factorization (dataset: MovieLens).
- Clustering analysis of Bay Area restaurants (dataset: Yelp)
- Social information based Recommender system (project used social connections to make restaurant recommendations to users. Datasets: Facebook and Yelp)
- Personalized news feed using Facebook data (datasets: Facebook and Yahoo! News)
- SMS classification (dataset: SMS Spam from UCI repository)
- Identification of spam/fraudulent advertisements on craigslist (dataset: craigslist)
- A Facebook-based movie recommendation system (dataset: Facebook and MovieLens)
- Song recommendation engine (dataset: Last.fm)
- Crime recognition and prediction system (dataset: SFPD crime incident database)
- Forecasting elections using sentiment analysis on tweets (dataset: Twitter).
- Polarity analysis of movie reviews (dataset: MovieLens)
- Book recommendation systems (dataset: BookCrossing)
- Analysis of rent prices given external factors (such as nearby schools, jails, etc.) (the students used 3 different datasets to perform their analysis (scraped data from zillow.com, used yelp, and census data).
- Event recommendation system (dataset: Facebook and and meetup data)
- Finding smartphone trends using sentiment analysis on tweets (dataset: Twitter).

Project report template

Section 1 Introduction

- Motivation
- Objective
- (Literature/Market review)

Section 2 System Design & Implementation details

- Algorithm(s) considered/selected (and why)
- Technologies & Tools used (and why)
- Architecture-related decisions (if applicable)
- System (and subsystems if needed) design/architecture/data flow (you may use diagrams with some supportive text)
- (Component details)
- Use cases / GUI / screenshots (as applicable)

Section 3 Experiments / Proof of concept evaluation

- Dataset(s) used (name, source, type of data, size of data, # of instances/statistics, any preprocessing performed etc.)
- Methodology followed (e.g. n-fold-cross validation, number of n, size of training/test set etc.) (as applicable)
- Graphs showing different parameters/algorithms evaluated in a comparative manner, along with some supportive text. (as applicable)
- Analysis of results

Section 4 Discussion & Conclusions (bullet points as applicable)

- Decisions made
- Difficulties faced
- Things that worked
- Things that didn't work well
- Conclusion

Section 5 Project Plan / Task Distribution

- Break into components and clearly explain
 - Who was assigned to what task
 - Who ended up doing what task (justify as applicable)

Project grading rubric

Criteria	Ratings				Pts	
Project report: completeness	The team has included all necessary information for all critical parts of their project, including enough details to allow for someone who is not familiar with the project to understand the entire process. 10.0 pts		The report includes sufficient information regarding the project work, but some important details are missing. 5.0 pts		Many important details are missing from the report. 0.0 pts	10.0 pts
Project report: template	Template was followed 100% and maximum number of pages was not exceeded. 4.0 pts	Template was followed for the most part, but some things are missing and/or over page limit. 2.0 pts		Many things are missing and the template has not been properly followed 0.0 pts		4.0 pts
Project report: clarity	The report is well written, does not have grammar or syntax errors, and the formatting is consistent throughout, making it easy to read. 2.0 pts		The report is written well for the most part, some errors exist, or some inconsistencies in terms of the formatting exist. 1.0 pts		The report contains typos, errors, and/or has formatting issues. 0.0 pts	2.0 pts
Prototype: dataset & problems addressed	The dataset used for this project is big enough or complex enough so that the project is within scope. The problems addressed are interesting and relevant. 6.0 pts	The dataset used for this project is sufficient but the team could have used alternatives that would have made the scope more appropriate. Some of the problems addressed could be solved without the use of data mining. 3.0 pts		The dataset used and/or the problems used is either inappropriate for the project, or too small/simple so that it narrows significantly the project's scope. 0.0 pts		6.0 pts
Prototype: preprocessing and algorithms used	The pre-processing (where applicable) and data mining / analysis algorithms (and variations) used for this project were appropriately selected. All selection considerations are documented in the report. 12.0 pts		The team used a subset of the pre-processing/data mining techniques that should have been used. 6.0 pts		Only a simple algorithm used and no justification is provided and/or the selection was not appropriate for the given problem. 0.0 pts	12.0 pts
Prototype: Evaluation	The various algorithms were appropriately evaluated and/or compared against other algorithms or variations, using appropriate techniques learnt in this course (additional evaluation techniques are accepted). 10.0 pts		Some basic evaluation has been performed, but not to a full extend. 5.0 pts		No experimental evaluation was performed. 0.0 pts	10.0 pts
Presentation slides	Slides following the template, appropriate information. 6.0 pts	Slides followed the template but some information was missing. 3.0 pts		Big deviation from the template and/or confusing slides. 0.0 pts		6.0 pts
Presentation (in-class)	Well-timed, lasted 8 or less minutes, all members participated equally. 5.0 pts	Run ~1 minute over. Or not equal member participation (those who take more time will be deducted points) 3.0 pts			Run >1 minute late. Unequal participation. 0.0 pts	5.0 pts
Total Points: 55.0						